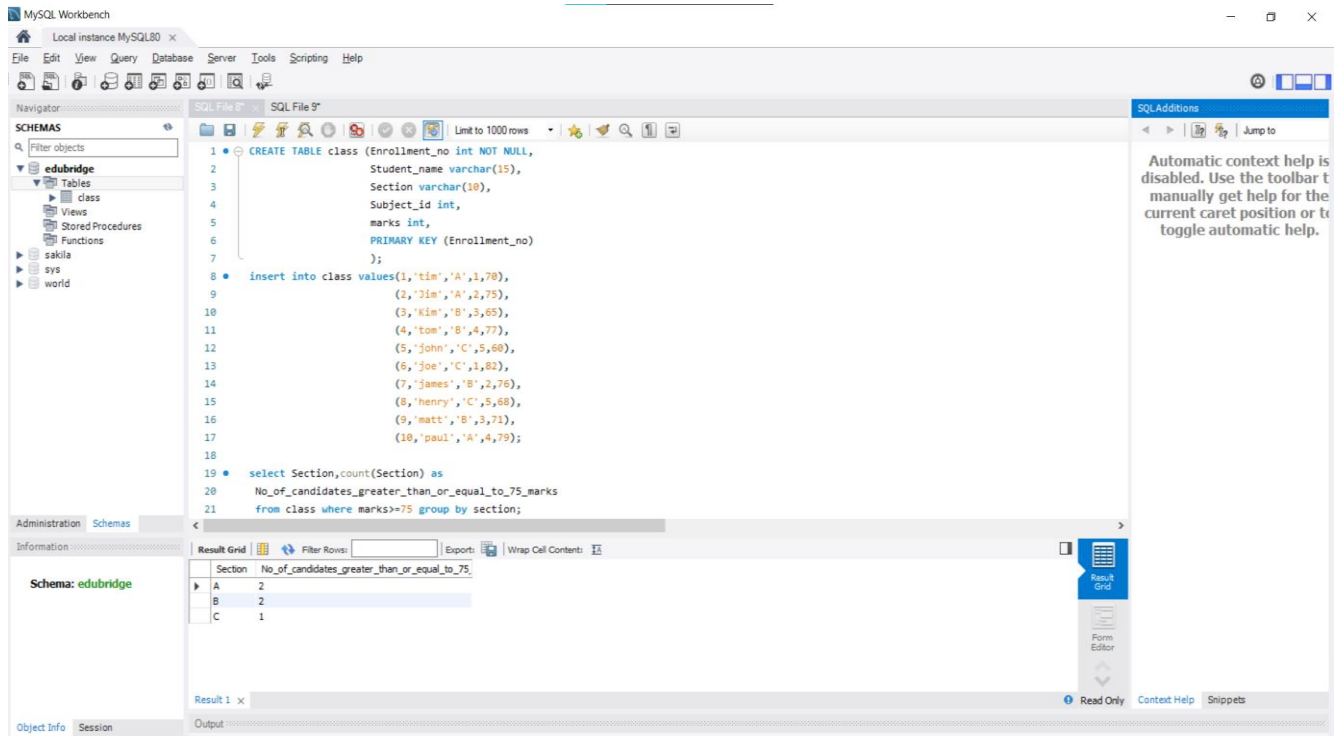


EDUBRIDGE ASSIGNMENT

Task -1:SQL

Answer:



```
CREATE TABLE class (Enrollment_no int NOT NULL,  
Student_name varchar(15),  
Section varchar(10),  
Subject_id int,  
marks int,  
PRIMARY KEY (Enrollment_no)  
);
```

```
insert into class values(1,'tim','A',1,70),  
                        (2,'Jim','A',2,75),  
                        (3,'Kim','B',3,65),  
                        (4,'tom','B',4,77),
```

```
(5,'john','C',5,60),
(6,'joe','C',1,82),
(7,'james','B',2,76),
(8,'henry','C',5,68),
(9,'matt','B',3,71),
(10,'paul','A',4,79);
```

```
select Section,count(Section) as
No_of_candidates_greater_than_or_equal_to_75_marks
from class where marks>=75 group by section;
```

output:

Result Grid

Filter Rows:

Export:

Wrap Cell Content:

	Section	No_of_candidates_greater_than_or_equal_to_75
▶	A	2
	B	2
	C	1

Result Grid

Form Editor

Task 2: Tableau

Answer:

Load data from excel and select rows

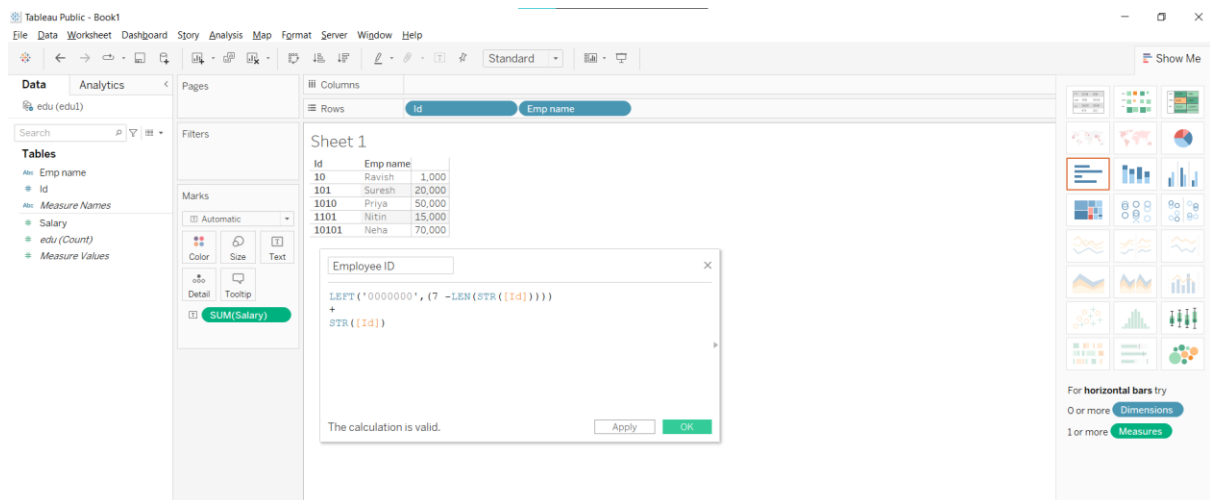
The screenshot displays the Tableau Public interface for a worksheet named "edu (edu1)". The top menu bar includes File, Data, Worksheet, Dashboard, Story, Analysis, Map, Format, Server, Window, and Help. The toolbar contains various icons for navigation and analysis. The sidebar on the right features a "Show Me" button. The main workspace is divided into four panes: Pages, Columns, Rows, and Marks. The Columns shelf contains "Id" and "Emp name". The Rows shelf is empty. The Marks shelf is set to "Automatic". The data source is "edu (edu1)". The data table shows employee details with columns "Id", "Emp name", and "Salary". The "Salary" field is highlighted in the Marks shelf, and a green pill labeled "SUM(Salary)" is visible in the Marks card.

Id	Emp name	Salary
10	Ravish	1,000
101	Suresh	20,000
1010	Priya	50,000
1101	Nitin	15,000
10101	Neha	70,000

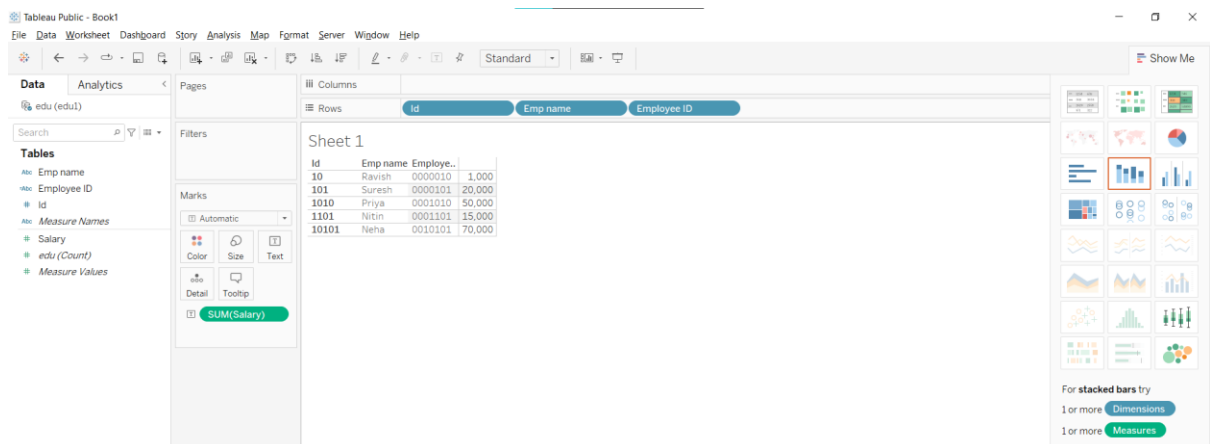
Then create calculated field to achieve the given condition using the formula

`LEFT('0000000',(7 -LEN(STR([Id]))))+STR([Id])`

And name the row as Employee Id as show below

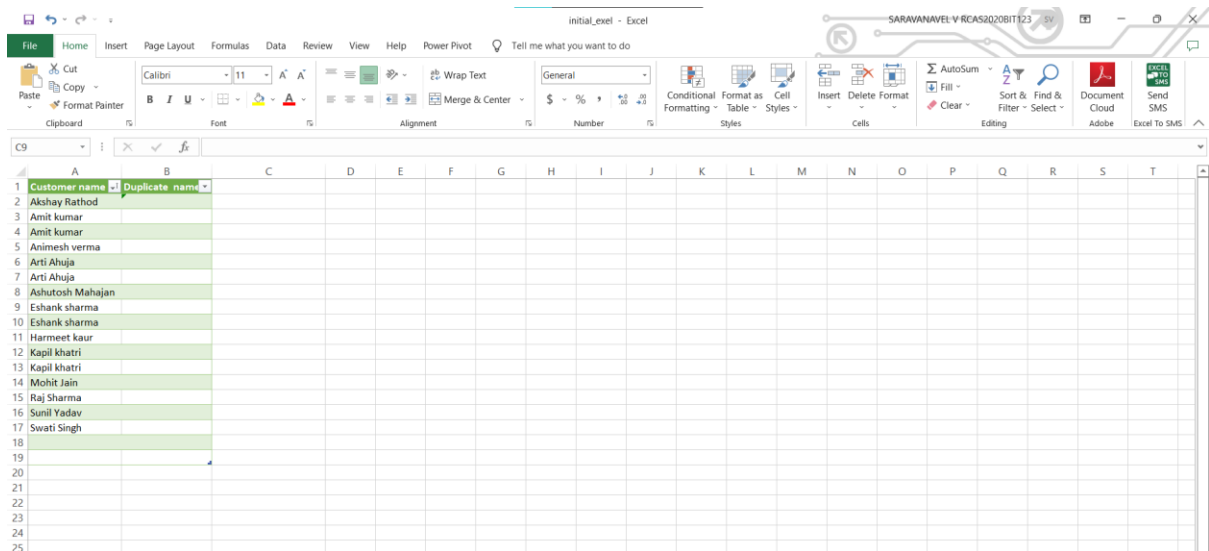


After applying the formula we can achieve the task as shown below

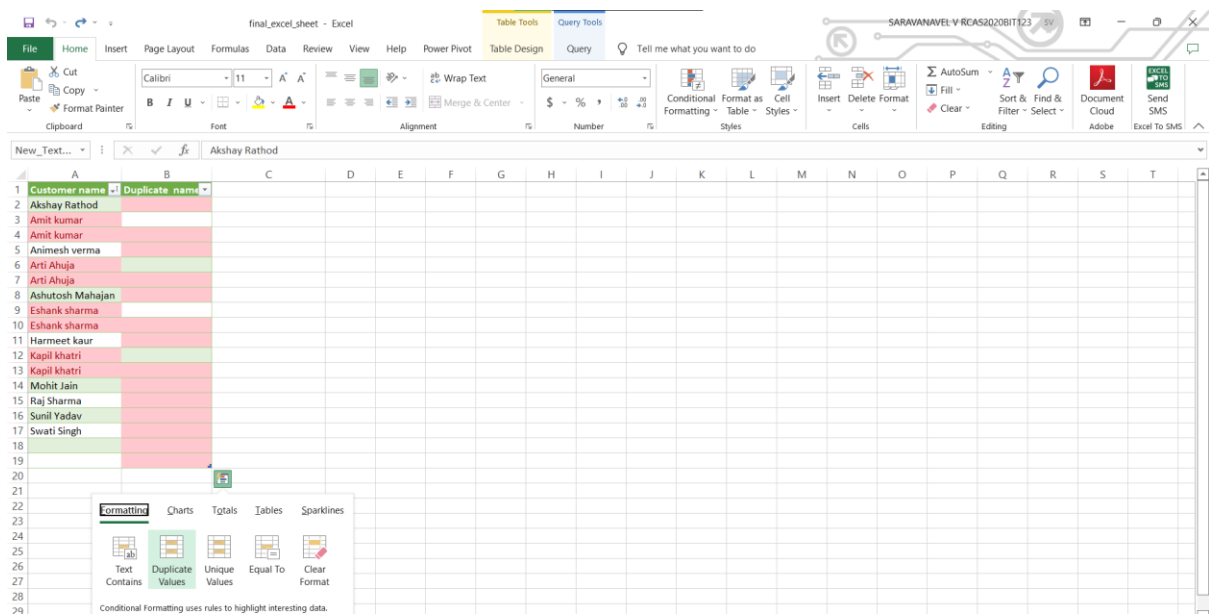


Task 3:EXCEL

Load the given table in excel



Then apply conditional formatting for duplicate values as shown below



Then apply the given formula and drag it to all cells

=IFERROR(INDEX(A3, MATCH(0, COUNTIF(B\$1, A3:A12)+IF(COUNTIF(A3:A12, A3:A12)>1, 0, 1), 0)), "")

final_excel_sheet - Excel

Table Tools

File Home Insert Page Layout Formulas Data Review View Help Power Pivot Table Design Tell me what you want to do

Clipboard: Cut, Copy, Paste, Format Painter

Font: Calibri, 11, Bold, Italic, Underline, Text Color, Background Color

Alignment: Wrap Text, Merge & Center

Number: General, Currency, Percentage, Date, Time, Text, Fraction, Scientific

Conditional Formatting: Conditional Formatting, Format as Table, Cell Styles

Cells: Insert, Delete, Format

Editing: AutoSum, Fill, Sort & Filter, Find & Select, Document Cloud, Send SMS

Formula Bar: B3 {=IFERROR(INDEX(A3, MATCH(0, COUNTIF(B\$1, A3:A12)+IF(COUNTIF(A3:A12, A3:A12)>1, 0, 1), 0)), "")}

Customer name	Duplicate name
Akhay Rathod	
Amit kumar	Amit kumar
Amit kumar	
Animesh verma	
Arti Ahuja	Arti Ahuja
Arti Ahuja	
Ashutosh Mahajan	
Eshank sharma	Eshank sharma
Eshank sharma	
Harmeet kaur	
Kapil khatri	Kapil khatri
Kapil khatri	
Mohit Jain	
Raj Sharma	
Sunil Yadav	
Swati Singh	

Task 4:Machine Learning

Machine Learning Task

classify fetal health to avoid the abnormalities for the child and mother while giving birth

Importing libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
#from sklearn.svm import LinearSVC
#from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report
from sklearn import metrics
#from sklearn.metrics import roc_curve, auc, roc_auc_score
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
```

Load Data

```
In [3]: train_data = pd.read_csv("train.csv")
test_data=pd.read_csv("test.csv")
val_data=pd.read_csv("val.csv")
```

comparing test and validation data

```
In [4]: test_data.shape
```

```
Out[4]: (426, 21)
```

```
In [5]: val_data.shape
```

```
Out[5]: (426, 21)
```

```
In [6]: com=test_data[test_data.apply(tuple,1).isin(val_data.apply(tuple,1))]
```

```
In [7]: scv=com.to_csv("sample.csv",index=False)
```

```
In [8]: diff=pd.read_csv("sample.csv")
```

```
In [9]: diff.shape
```

```
Out[9]: (426, 21)
```

As seen there is no difference between test and validation set

Pre-processing and EDA

```
In [10]: train_data.head()
```

```
Out[10]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolong
0	142.0	0.000	0.000	0.007	0.000	0.0	
1	122.0	0.000	0.000	0.006	0.002	0.0	
2	129.0	0.005	0.003	0.001	0.000	0.0	
3	136.0	0.006	0.000	0.008	0.000	0.0	
4	144.0	0.000	0.000	0.006	0.000	0.0	

5 rows × 22 columns

```
In [11]: train_data.isnull().sum() #No null values
```

```
Out[11]:
```

baseline value	0
accelerations	0
fetal_movement	0
uterine_contractions	0
light_decelerations	0
severe_decelerations	0
prolongued_decelerations	0
abnormal_short_term_variability	0
mean_value_of_short_term_variability	0
percentage_of_time_with_abnormal_long_term_variability	0
mean_value_of_long_term_variability	0
histogram_width	0
histogram_min	0
histogram_max	0
histogram_number_of_peaks	0
histogram_number_of_zeroes	0
histogram_mode	0
histogram_mean	0
histogram_median	0
histogram_variance	0
histogram_tendency	0
fetal_health	0
dtype: int64	

```
In [12]: train_data.describe()
```

```
Out[12]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations
count	1700.000000	1700.000000	1700.000000	1700.000000	1700.000000	1700.000000
mean	133.213529	0.003212	0.010211	0.004356	0.001899	0.000004
std	9.873344	0.003888	0.050124	0.002943	0.002976	0.000059
min	106.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	126.000000	0.000000	0.000000	0.002000	0.000000	0.000000
50%	133.000000	0.002000	0.000000	0.004000	0.000000	0.000000

75%	140.000000	0.006000	0.003000	0.006000	0.003000	0.000000
max	159.000000	0.019000	0.481000	0.015000	0.015000	0.001000

8 rows × 22 columns

```
In [13]: train_data.duplicated().sum()
```

```
Out[13]: 7
```

```
In [14]: train_df=train_data.drop_duplicates()
```

```
In [15]: train_df.duplicated().sum()
```

```
Out[15]: 0
```

```
In [16]: train_df.shape
```

```
Out[16]: (1693, 22)
```

```
In [17]: train_df.nunique()
```

```
Out[17]: baseline value                47
accelerations                20
fetal_movement               92
uterine_contractions         16
light_decelerations          16
severe_decelerations          2
prolongued_decelerations      6
abnormal_short_term_variability 75
mean_value_of_short_term_variability 57
percentage_of_time_with_abnormal_long_term_variability 84
mean_value_of_long_term_variability 240
histogram_width              153
histogram_min                 109
histogram_max                 85
histogram_number_of_peaks     18
histogram_number_of_zeroes     7
histogram_mode                84
histogram_mean                99
histogram_median              91
histogram_variance            127
histogram_tendency            3
fetal_health                  3
dtype: int64
```


```
In [18]: train_df['fetal_health'].unique()
```

```
Out[18]: array([1., 3., 2.])
```


```
In [19]: train_df['fetal_health'].value_counts()
```

```
Out[19]: 1.0    1317
2.0     235
3.0     141
Name: fetal_health, dtype: int64
```

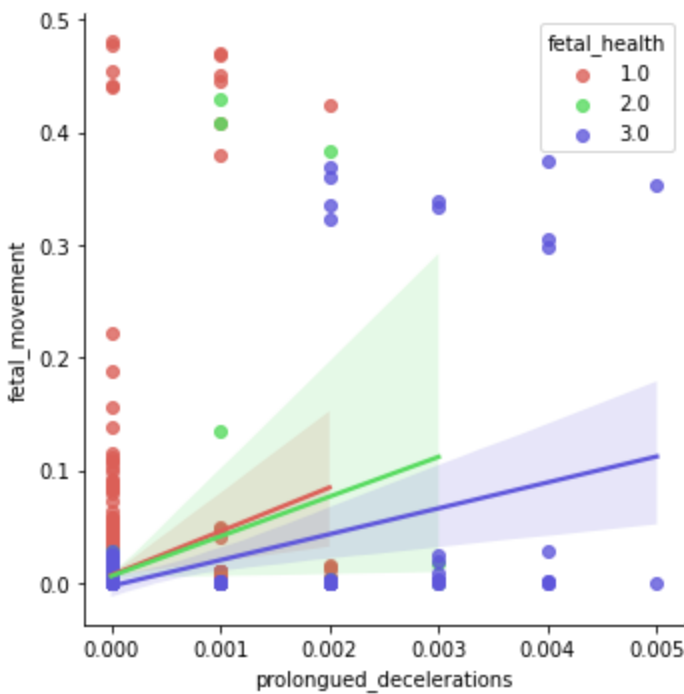
```
In [20]: plt.figure(figsize=(15,15))
cmap = sns.diverging_palette(250,10, s=80, l=55, n=9, as_cmap=True)
sns.heatmap(train_df.corr(),annot=True,cmap = cmap, center=0)
plt.show()
```


```
In [21]: sns.lmplot(data = train_df, x = "accelerations", y = "fetal_movement", palette = 'hls',
plt.show())
```




```
In [22]: sns.lmplot(data = train_df, x = "prolonged_decelerations", y = "fetal_movement", palette="magma",
plt.show())
```



```
In [23]: sns.lmplot(data = train_df, x = "abnormal_short_term_variability", y = "fetal_movement",
plt.show())
```



```
In [24]: sns.lmplot(data = train_df, x = "mean_value_of_long_term_variability", y = "fetal_movement",
plt.show())
```



Model building

```
In [25]: x = train_df.drop(['fetal_health'],axis=1)
y = train_df['fetal_health']
```

```
In [26]: print(x.shape)
print(y.shape)
```

```
(1693, 21)
(1693,)
```

```
In [27]: column_names = list(x.columns)
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_df = pd.DataFrame(x_scaled, columns = column_names)
```

```
In [28]: x_train, x_test, y_train, y_test = train_test_split(x_df, y, test_size = 0.2, random_state=42)

In [29]: print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

(1354, 21)
(1354,)
(339, 21)
(339,)

In [30]: pipeline_lr = Pipeline([('lr_classifier', LogisticRegression(random_state= 42))])
pipeline_dt = Pipeline([('dt_classifier', DecisionTreeClassifier(random_state= 42))])
pipeline_rf = Pipeline([('rf_classifier', RandomForestClassifier())])
pipeline_svc = Pipeline([('sv_classifier', SVC())])
pipelines = [pipeline_lr, pipeline_dt, pipeline_rf, pipeline_svc]
pipe_dict = {0:'Logistic Regression', 1:'Decision Tree', 2:'Random Forest', 3:'SVC'}

In [31]: for pipe in pipelines:
pipe.fit(x_train, y_train)
```

Cross validation score

```
In [32]: cv_results_accuracy = []
for i, model in enumerate(pipelines):
    cv_score = cross_val_score(model, x_train, y_train, cv = 10)
    cv_results_accuracy.append(cv_score)
    print("%s: %f " % (pipe_dict[i], cv_score.mean()))

Logistic Regression: 0.880381
Decision Tree: 0.923932
Random Forest: 0.932048
SVC: 0.892941
```

As seen above the cross validation score of random forest is high.so selecting random forest for model prediction

```
In [33]: pred_rfc = pipeline_rf.predict(x_test)

In [34]: print("Classification Report for classifier %s:\n%s\n" % (pipeline_rf, metrics.classification_report(y_test, pred_rfc)))

Classification Report for classifier Pipeline(steps=[('rf_classifier', RandomForestClassifier())]):
              precision    recall  f1-score   support

     1.0         0.93      0.99      0.96         263
     2.0         0.93      0.58      0.72          48
     3.0         0.89      0.86      0.87          28

 accuracy          0.92
 macro avg         0.92
 weighted avg      0.92
```


confusion matrix and accuracy

```
In [35]: matrix = confusion_matrix(y_test, pred_rfc)
print('Confusion Matrix : \n', matrix)
```

Confusion Matrix :

```
[[261  1  1]
 [ 18 28  2]
 [  3  1 24]]
```

```
In [36]: from sklearn.metrics import plot_confusion_matrix
plot_confusion_matrix(pipeline_rf, x_test, y_test, cmap="plasma")
plt.show()
```



```
In [37]: print("Accuracy"+str(accuracy_score(y_test,pred_rfc)))
Accuracy0.9233038348082596
```

Saving Model prediction

```
In [38]: pred_rfc
```

```
Out[38]: array([1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1.,
        1., 3., 1., 1., 1., 2., 1., 1., 3., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 3., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 3., 3., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 1., 1.,
        3., 1., 1., 1., 1., 1., 1., 1., 3., 1., 1., 1., 1., 2., 1., 2., 1.,
        1., 2., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1.,
        1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 2., 1., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 2., 1., 2., 1.,
        1., 1., 1., 1., 1., 1., 1., 2., 1., 3., 1., 1., 1., 3., 1., 3., 3.,
        1., 2., 1., 1., 3., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 2.,
        1., 1., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 2., 1., 1., 1., 1., 3., 2., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 2., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 1., 3.,
        1., 1., 1., 1., 3., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 3., 1., 1., 2., 1., 1., 1., 2., 1., 1., 1., 1., 3.,
        1., 1., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 3., 1., 1.,
        1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 2., 1., 1., 3., 1., 1.]])
```

```
In [39]: df = pd.DataFrame(pred_rfc, columns =['fetal_health'])
```

```
In [40]: df.head()
```

```
Out[40]:
```

	fetal_health
0	1.0
1	1.0

2	1.0
3	2.0
4	1.0

```
In [41]: df.to_csv('output.csv',header=df.columns,index=False)
```

```
In [42]: dfg=pd.read_csv("output.csv")
dfg.head()
```

```
Out[42]:
```

	fetal_health
0	1.0
1	1.0
2	1.0
3	2.0
4	1.0

```
In [43]: dfg.shape
```

```
Out[43]: (339, 1)
```