

Machine Learning Task

classify fetal health to avoid the abnormalities for the child and mother while giving birth

Importing libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
#from sklearn.svm import LinearSVC
#from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report
from sklearn import metrics
#from sklearn.metrics import roc_curve, auc, roc_auc_score
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
```

Load Data

```
In [3]: train_data = pd.read_csv("train.csv")
test_data=pd.read_csv("test.csv")
val_data=pd.read_csv("val.csv")
```

comparing test and validation data

```
In [4]: test_data.shape
```

```
Out[4]: (426, 21)
```

```
In [5]: val_data.shape
```

```
Out[5]: (426, 21)
```

```
In [6]: com=test_data[test_data.apply(tuple,1).isin(val_data.apply(tuple,1))]
```

```
In [7]: scv=com.to_csv("sample.csv",index=False)
```

```
In [8]: diff=pd.read_csv("sample.csv")
```

```
In [9]: diff.shape
```

```
Out[9]: (426, 21)
```

As seen there is no difference between test and validation set

Pre-processing and EDA

```
In [10]: train_data.head()
```

```
Out[10]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations	prolong
0	142.0	0.000	0.000	0.007	0.000	0.0	
1	122.0	0.000	0.000	0.006	0.002	0.0	
2	129.0	0.005	0.003	0.001	0.000	0.0	
3	136.0	0.006	0.000	0.008	0.000	0.0	
4	144.0	0.000	0.000	0.006	0.000	0.0	

5 rows × 22 columns

```
In [11]: train_data.isnull().sum() #No null values
```

```
Out[11]:
```

baseline value	0
accelerations	0
fetal_movement	0
uterine_contractions	0
light_decelerations	0
severe_decelerations	0
prolongued_decelerations	0
abnormal_short_term_variability	0
mean_value_of_short_term_variability	0
percentage_of_time_with_abnormal_long_term_variability	0
mean_value_of_long_term_variability	0
histogram_width	0
histogram_min	0
histogram_max	0
histogram_number_of_peaks	0
histogram_number_of_zeroes	0
histogram_mode	0
histogram_mean	0
histogram_median	0
histogram_variance	0
histogram_tendency	0
fetal_health	0
dtype: int64	

```
In [12]: train_data.describe()
```

```
Out[12]:
```

	baseline value	accelerations	fetal_movement	uterine_contractions	light_decelerations	severe_decelerations
count	1700.000000	1700.000000	1700.000000	1700.000000	1700.000000	1700.000000
mean	133.213529	0.003212	0.010211	0.004356	0.001899	0.000004
std	9.873344	0.003888	0.050124	0.002943	0.002976	0.000059
min	106.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	126.000000	0.000000	0.000000	0.002000	0.000000	0.000000
50%	133.000000	0.002000	0.000000	0.004000	0.000000	0.000000

75%	140.000000	0.006000	0.003000	0.006000	0.003000	0.000000
max	159.000000	0.019000	0.481000	0.015000	0.015000	0.001000

8 rows × 22 columns

```
In [13]: train_data.duplicated().sum()
```

```
Out[13]: 7
```

```
In [14]: train_df=train_data.drop_duplicates()
```

```
In [15]: train_df.duplicated().sum()
```

```
Out[15]: 0
```

```
In [16]: train_df.shape
```

```
Out[16]: (1693, 22)
```

```
In [17]: train_df.nunique()
```

```
Out[17]: baseline value                47
accelerations                        20
fetal_movement                      92
uterine_contractions                 16
light_decelerations                 16
severe_decelerations                 2
prolongued_decelerations             6
abnormal_short_term_variability      75
mean_value_of_short_term_variability 57
percentage_of_time_with_abnormal_long_term_variability 84
mean_value_of_long_term_variability  240
histogram_width                      153
histogram_min                        109
histogram_max                        85
histogram_number_of_peaks            18
histogram_number_of_zeroes           7
histogram_mode                       84
histogram_mean                       99
histogram_median                     91
histogram_variance                   127
histogram_tendency                   3
fetal_health                         3
dtype: int64
```


```
In [18]: train_df['fetal_health'].unique()
```

```
Out[18]: array([1., 3., 2.])
```


```
In [19]: train_df['fetal_health'].value_counts()
```

```
Out[19]: 1.0    1317
2.0     235
3.0     141
Name: fetal_health, dtype: int64
```


```
In [20]: plt.figure(figsize=(15,15))
cmap = sns.diverging_palette(250,10, s=80, l=55, n=9, as_cmap=True)
sns.heatmap(train_df.corr(),annot=True,cmap = cmap, center=0)
plt.show()
```



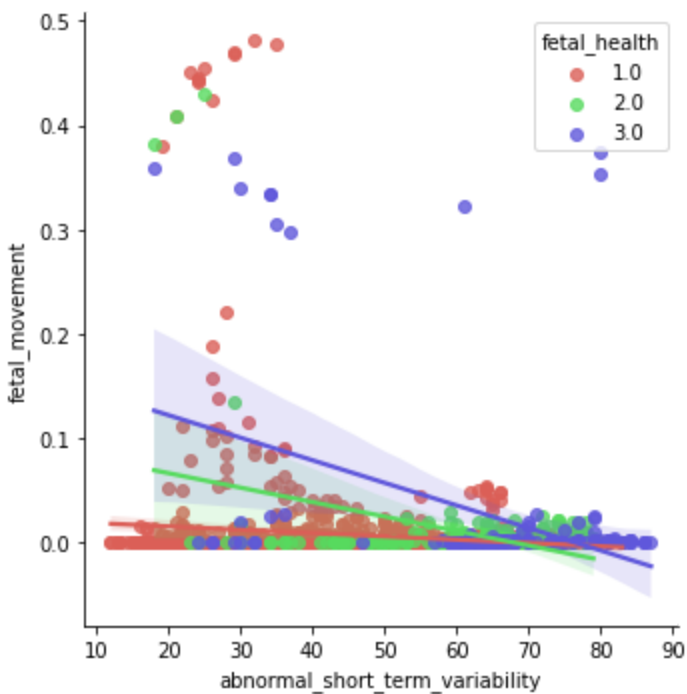
```
In [21]: sns.lmplot(data = train_df, x = "accelerations", y = "fetal_movement", palette = 'hls',
plt.show())
```




```
In [22]: sns.lmplot(data = train_df, x = "prolongued_decelerations", y = "fetal_movement", palette="magma",
plt.show())
```



```
In [23]: sns.lmplot(data = train_df, x = "abnormal_short_term_variability", y = "fetal_movement",
plt.show())
```



```
In [24]: sns.lmplot(data = train_df, x = "mean_value_of_long_term_variability", y = "fetal_movement",
plt.show())
```



Model building

```
In [25]: x = train_df.drop(['fetal_health'],axis=1)
y = train_df['fetal_health']
```

```
In [26]: print(x.shape)
print(y.shape)
```

```
(1693, 21)
(1693,)
```

```
In [27]: column_names = list(x.columns)
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_df = pd.DataFrame(x_scaled, columns = column_names)
```

```
In [28]: x_train, x_test, y_train, y_test = train_test_split(x_df, y, test_size = 0.2, random_state=42)

In [29]: print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)

(1354, 21)
(1354,)
(339, 21)
(339,)

In [30]: pipeline_lr = Pipeline([('lr_classifier', LogisticRegression(random_state= 42))])
pipeline_dt = Pipeline([('dt_classifier', DecisionTreeClassifier(random_state= 42))])
pipeline_rf = Pipeline([('rf_classifier', RandomForestClassifier())])
pipeline_svc = Pipeline([('sv_classifier', SVC())])
pipelines = [pipeline_lr, pipeline_dt, pipeline_rf, pipeline_svc]
pipe_dict = {0:'Logistic Regression', 1:'Decision Tree', 2:'Random Forest', 3:'SVC'}

In [31]: for pipe in pipelines:
pipe.fit(x_train, y_train)
```

Cross validation score

```
In [32]: cv_results_accuracy = []
for i, model in enumerate(pipelines):
    cv_score = cross_val_score(model, x_train, y_train, cv = 10)
    cv_results_accuracy.append(cv_score)
    print("%s: %f " % (pipe_dict[i], cv_score.mean()))

Logistic Regression: 0.880381
Decision Tree: 0.923932
Random Forest: 0.932048
SVC: 0.892941
```

As seen above the cross validation score of random forest is high.so selecting random forest for model prediction

```
In [33]: pred_rfc = pipeline_rf.predict(x_test)

In [34]: print("Classification Report for classifier %s:\n%s\n" % (pipeline_rf, metrics.classification_report(y_test, pred_rfc)))

Classification Report for classifier Pipeline(steps=[('rf_classifier', RandomForestClassifier())]):

```

	precision	recall	f1-score	support
1.0	0.93	0.99	0.96	263
2.0	0.93	0.58	0.72	48
3.0	0.89	0.86	0.87	28
accuracy			0.92	339
macro avg	0.92	0.81	0.85	339
weighted avg	0.92	0.92	0.92	339


confusion matrix and accuracy

```
In [35]: matrix = confusion_matrix(y_test, pred_rfc)
print('Confusion Matrix : \n', matrix)
```

Confusion Matrix :

```
[[261  1   1]
 [ 18 28   2]
 [  3  1 24]]
```

```
In [36]: from sklearn.metrics import plot_confusion_matrix
plot_confusion_matrix(pipeline_rf, x_test, y_test, cmap="plasma")
plt.show()
```



```
In [37]: print("Accuracy"+str(accuracy_score(y_test,pred_rfc)))
Accuracy0.9233038348082596
```

Saving Model prediction

```
In [38]: pred_rfc
```

```
Out[38]: array([1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1.,
        1., 3., 1., 1., 1., 2., 1., 1., 3., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 3., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 2., 3., 3., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 1., 1.,
        3., 1., 1., 1., 1., 1., 1., 1., 3., 1., 1., 1., 1., 2., 1., 2., 1.,
        1., 2., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1.,
        1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 2., 1., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 2., 1., 2., 1.,
        1., 1., 1., 1., 1., 1., 1., 2., 1., 3., 1., 1., 1., 3., 1., 3., 3.,
        1., 2., 1., 1., 3., 1., 1., 1., 1., 1., 2., 1., 1., 1., 1., 1., 1., 2.,
        1., 1., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 2., 1., 1., 1., 1., 3., 2., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 2., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 3.,
        1., 1., 1., 3., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 3., 1., 1., 2., 1., 1., 1., 2., 1., 1., 1., 1., 1., 3.,
        1., 1., 1., 3., 1., 1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 3., 1., 1.,
        1., 1., 1., 1., 1., 1., 1., 2., 1., 1., 1., 2., 1., 1., 3., 1., 1.]])
```

```
In [39]: df = pd.DataFrame(pred_rfc, columns =['fetal_health'])
```

```
In [40]: df.head()
```

```
Out[40]:
```

	fetal_health
0	1.0
1	1.0

2	1.0
3	2.0
4	1.0

```
In [41]: df.to_csv('output.csv',header=df.columns,index=False)
```

```
In [42]: dfg=pd.read_csv("output.csv")
dfg.head()
```

```
Out[42]:
```

	fetal_health
0	1.0
1	1.0
2	1.0
3	2.0
4	1.0

```
In [43]: dfg.shape
```

```
Out[43]: (339, 1)
```